

HBase 簡介

軟體發展組

莊家雋

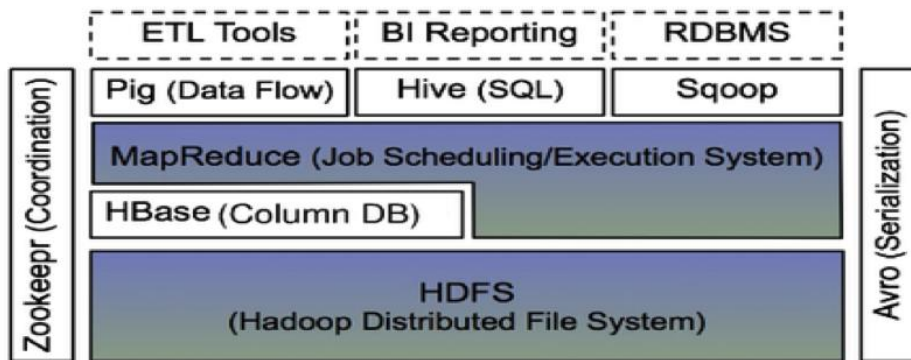
Outline

- What is Hbase
- Hbase Architecture
 - Master, region server, zookeeper
- Hbase Data model
- Hbase basic operations
 - Put, Get, Scan, Delete
- Hbase limit
- Performance issue
 - Rowkey, filter
- Hbase use case

What is HBase

- Hbase是一個高可靠性、高性能、column-orient、scalability 的分散式儲存系統

The Hadoop Ecosystem

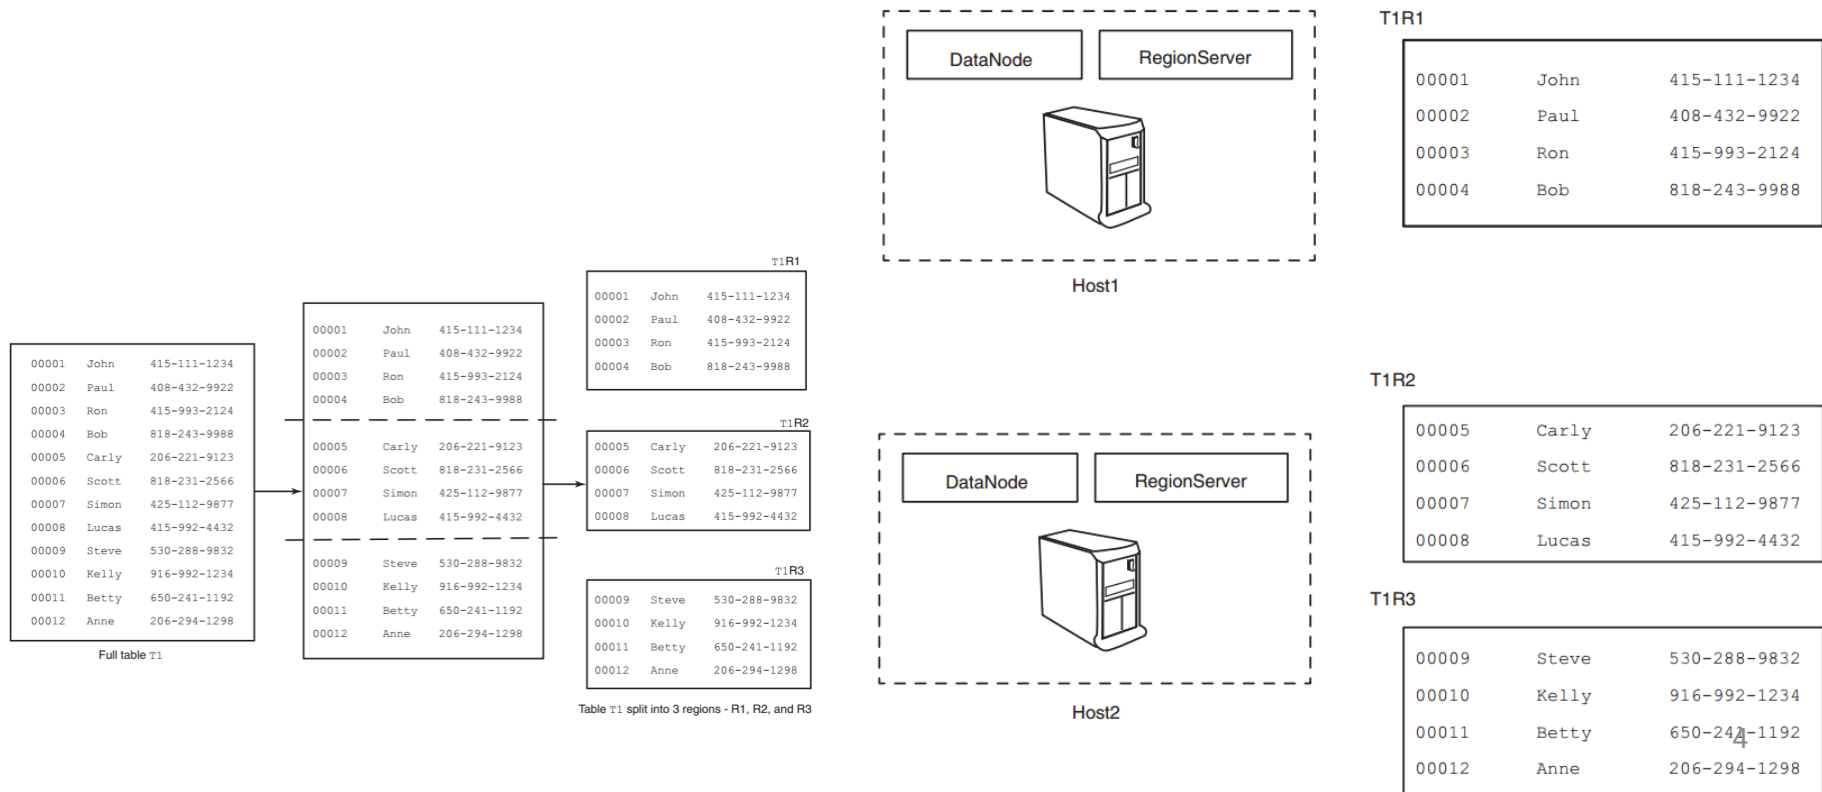


cloudera

Google	OpenSource
GFS	HDFS
MapReduce	Hadoop MapReduce
BigTable	HBase
Chubby	Zookeeper

Hbase architecture


- RegionServer
 - Table can be split into many region
 - Each RegionServer contains many regions
 - Add RS to horizontal scale up



Hbase architecture


- HMaster
 - Responsible for assigning regions to RegionServer
- Zookeeper
 - Use to find two special table –ROOT- and .META.

Step 1




Client -> ZooKeeper: Where's -ROOT-?

Step 2




ZooKeeper -> Client : It's at RegionServer RS1.

Step 3




Client -> -ROOT- table on RS1:
Which .META. region can tell me about
row 00009 from table T1?

Step 4



-ROOT- table on RS1 -> Client : .META. region M2
on RegionServer RS3 has that info.

Step 5



Client -> .META. region M2 on RS3: I want to read row 00009 from table T1. Which region can I find it in, and what RegionServer is serving it?

Step 6



.META. region M2 on RS3 -> Client: region T1R3 on RegionServer RS3.

Step 7




Client -> Region T1R3 on RS3: I want to read row 00009.

Step 8



Region T1R3 on RS3 -> Client: Ah, here you go.

Hbase architecture



Hbase data model

- Table
 - Hbase organize data into tables

Table					

Hbase data model

- Row and rowkey
 - Data is stored to its row
 - Rows are identified uniquely by their rowkey
 - Rowkeys are stored lexicographically
 - Rowkeys are always treated as byte[]

Table					
Rowkey					
R1					
R2					

Hbase data model

- Column family
 - Data within a row is grouped by column family (CF)
 - CF must be declared with table creation
 - CF cannot be add or delete
 - CF names are treated as String

Table					
Rowkey	CF1		CF2	CF3	
R1					
R2					

Hbase data model

- Column qualifier
 - Qualifier is used to address data
 - Qualifier need NOT be specified in advanced
 - Qualifier name is treated as byte[]
 - CF + qualifier can be seen as column in RDBMS
 - Represent by CF:qualifier

Table					
Rowkey	CF1		CF2	CF3	
	q1	q2	q1	q3	q4
R1					
R2					

Hbase data model

- Cell
 - Combination of rowkey, CF, qualifier uniquely identifies a cell
 - Data is stored in a cell, call value
 - Value is treated as byte[]

Table					
Rowkey	CF1		CF2	CF3	
	q1	q2	q1	q3	q4
R1	V1	v2			
R2	v1	v2	v3	v4	v5

Hbase data model

- Version
 - Values within a cell are versioned.
 - Versions are identified by timestamp, treated as long

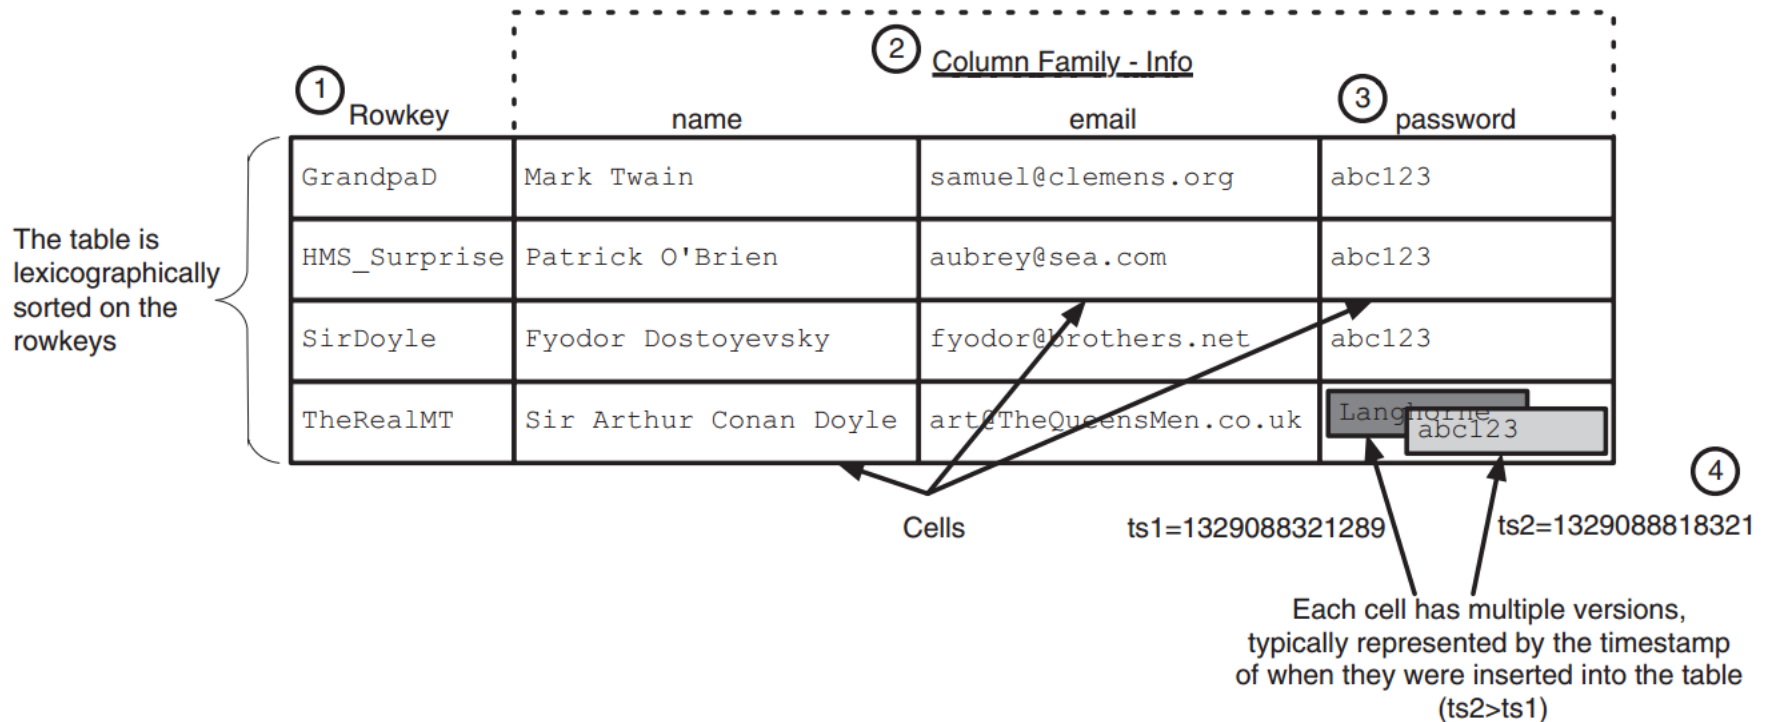
Table					
Rowkey	CF1		CF2	CF3	
	q1	q2	q1	q3	q4
R1	V1	v2			
R2	v1	v2	v3	v4	V5-1
					V5-2

Ver:1329088321289

Ver: 132908818321

Hbase data model

- A example put all together



Hbase basic operations

- Data manipulate
 - Put, Get, Scan
- Table administrate
 - Table Create 、 Disable 、 Drop
- Use Hbase shell and Java language

Hbase Shell

- `$ hbase shell # start hbase shell`
- In Hbase(main):
 - `create 't1','info'`
 - `put 't1','GrandpaD','info:name', 'mark Twain'`
 - `put 't1','GrandpaD','info:email','samuel@clemens.org'`
 - `put 't1','GrandpaD','info:password', 'ABC456'`
 - `put 't1','GrandpaD','info:password', 'abc123'`
 - `put 't1','GrandpaD','INFO:password', 'abc123' FAIL`
- `get 't1','GrandpaD'`
- `get 't1', 'GrandpaD', {COLUMN => 'info:password'}`
- `get 't1', 'GrandpaD', {COLUMN => 'info:password', VERSIONS => 3}`


- scan 't1'
- delete 't1', 'GrandpaD', 'info:password'
- deleteall 't1','GrandpaD'
- disable 't1'
- drop 't1'

Hbase Limit


- NO Table join
 - A single BIG table, Denormalize Design
- NO ACID
- NO secondary index
 - NO WHERE clause in SQL
- Row level transaction
 - Multiple rows modification is not thread safe
- C & P in CAP theorem

Performance issue

- Bad Rowkey design result in hotspot
 - Rowkey stored lexicographically
 - Avoid by Salting or hashing

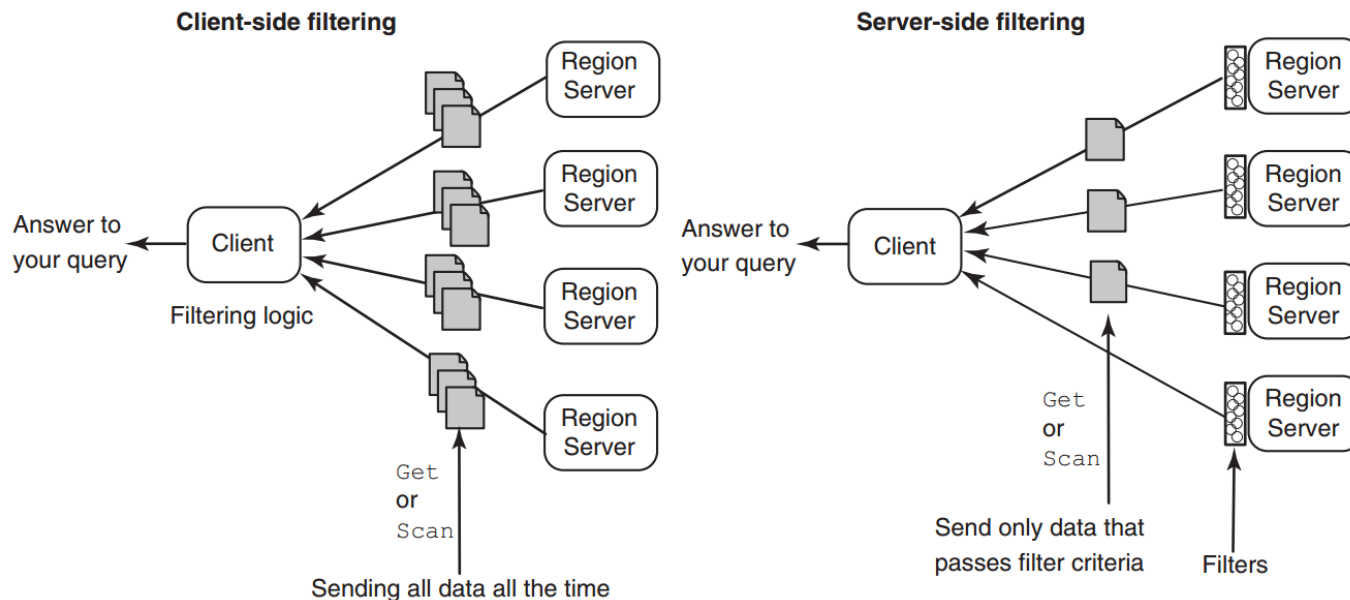


29



Performance issue

- Use Filter to reduce network/IO overhead



Hbase use case

- Count Statistics
- Store Graph

Statistics example

- By user
- By time unit
 - Per Day, per month , per year...
- By action type
 - Like, comment, ...

Initial Design

Table													
RK (userID)	TIME												
	20140901	20140902	...	20140931	20140900	20140901	...	20140923	201401	201402	...	201412	...
123	3			4									
987					15			9					

- Bad Design 1
 - Billion column per row is fine
 - Use column filter for query is BAD for performance
 - How about query by action ?


Initial Design

Table

Rowkey (userID)	Like				Dislike				Comment			
	2014	201401	...	20140914	2014	201401	...	20140914	2014	201401	...	20140914
123	3			4								
987					15			9				

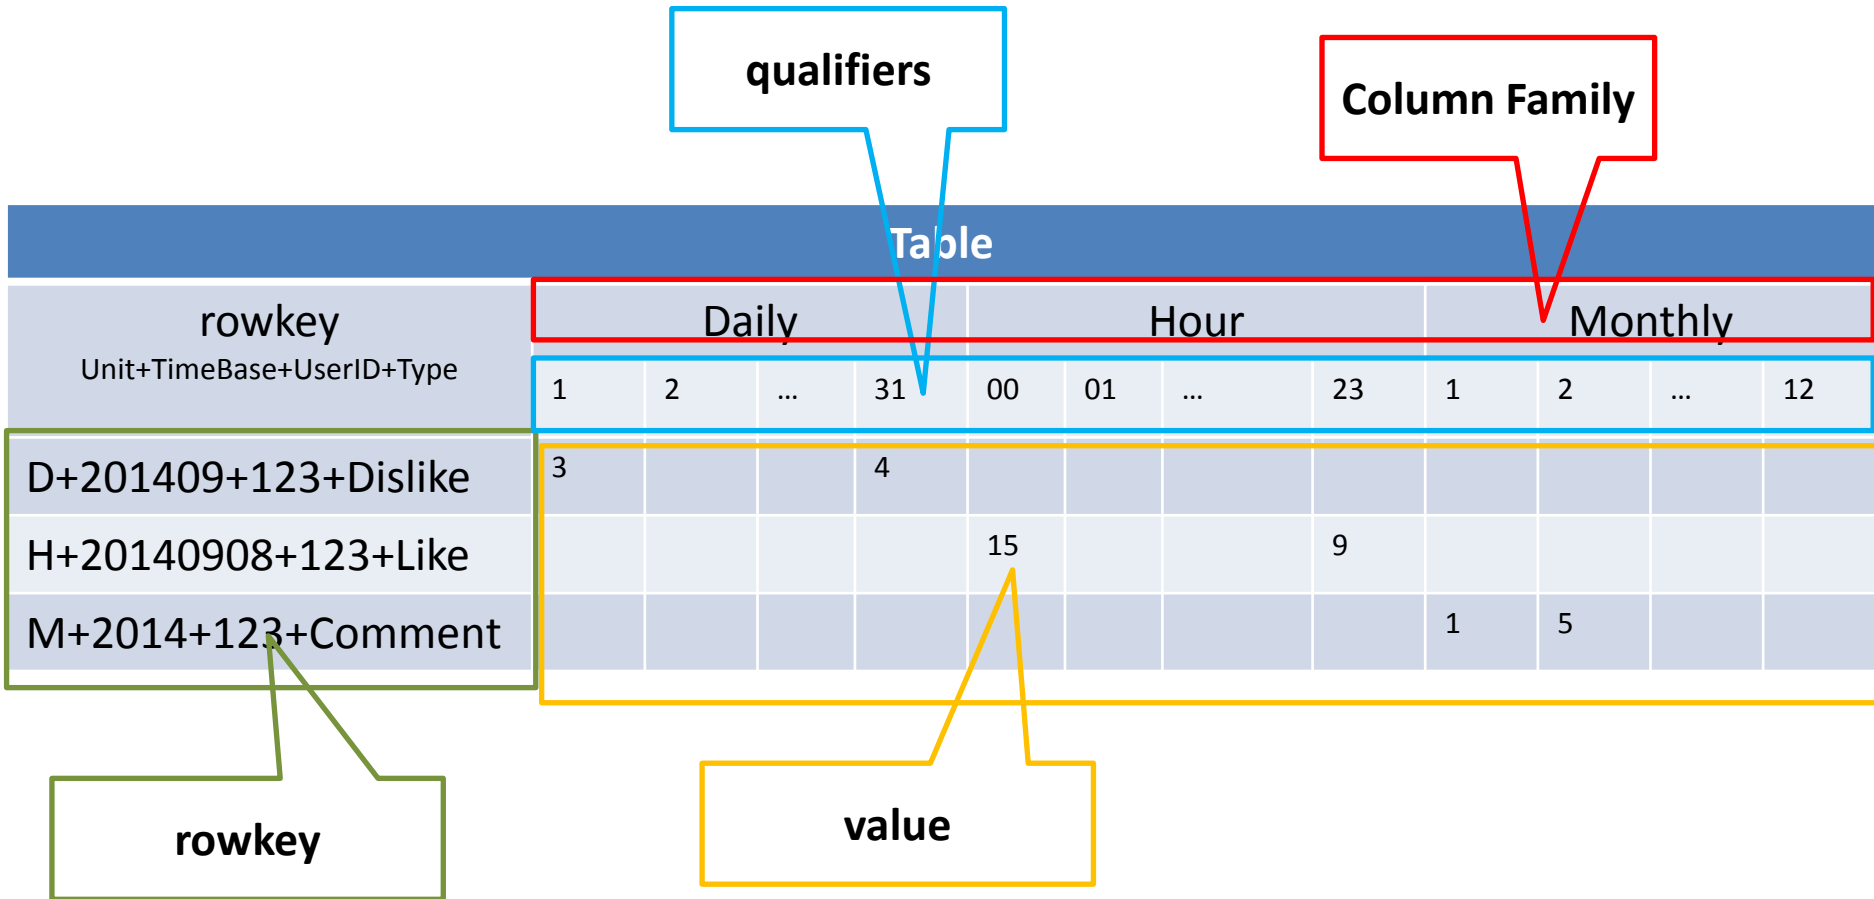
- Bad Design 2
 - CF should be defined first
 - How about add new action type ??
 - Number of CF should NOT be too much

Composition Rowkey



- Unit+TimeBase+UserID+Type
 - Unit : Char, (1 byte, H, D, M)
 - TimeBase: String
 - Length: 4 (Unit = M) or 6 (Unit = D) or 8 (Unit = H)
 - UserID: Long (8 bytes)
 - Type: Short (1 bytes)
 - 1 = Like, 2 = Dislike, 3 = comment

Better Design



主鍵查詢

rowkey query

- Get 789's Like action counts from from 2014/9/7 to 2014/9/20
 - Full RK: D + 201409 + DEF + Like

Table										
Rowkey (Unit+TimeBase+UserID+Type)	...	Daily								...
		...	7	8	9	...	19	20	...	
D+201409+123+Dislike			21	14	56	...	21	47		
D+201409+123+Like			25	12	78	...	98	112		
D+201409+123+comment			27	21	57	...	31	34		
D+201409+789+Like			26	41	29	...	7	35		
H+20140908+123+Like										
M+2014+123+Comment										

部分主鍵查詢


Partial rowkey query

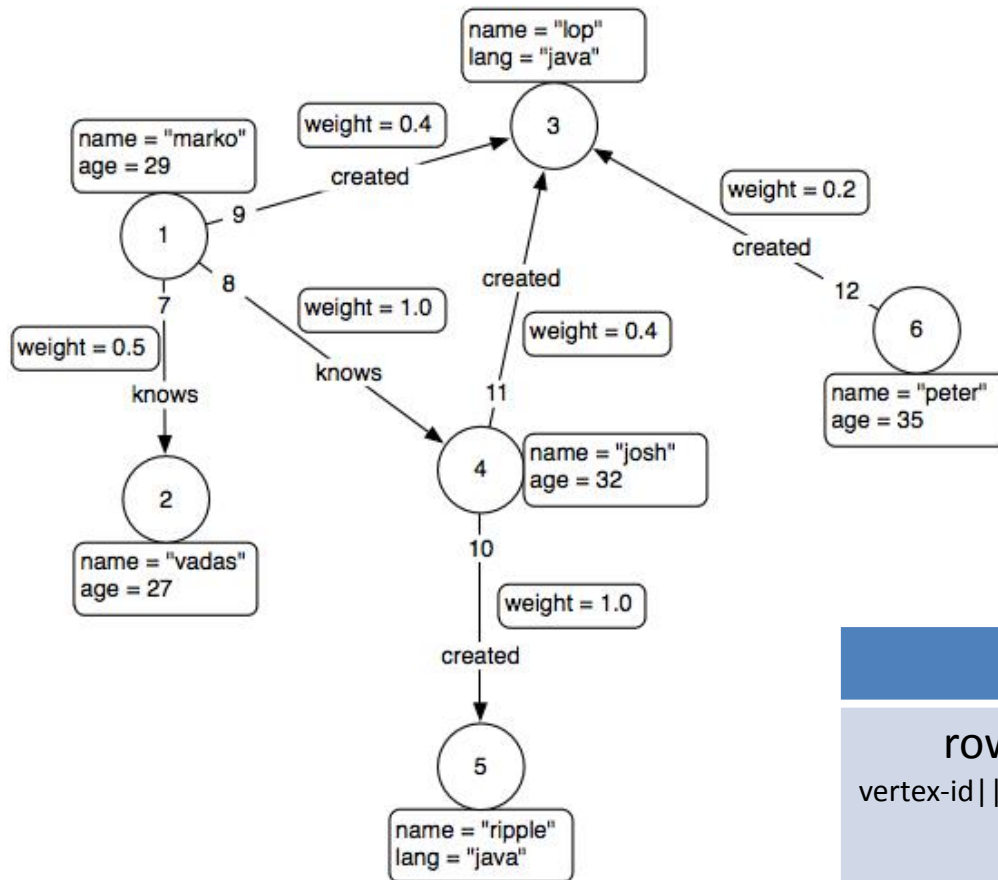
- Get 123's each action counts from from 2014/9/7 to 2014/9/20
 - Start RK: D+ 201409 + 123
 - END RK : D+ 201409 + 12**4**

Table										
Rowkey (Unit+TimeBase+UserID+Type)	...	Daily								...
		...	7	8	9	...	19	20	...	
D+201409+123+Dislike			21	14	56	...	21	47		
D+201409+123+Like			25	12	78	...	98	112		
D+201409+123+comment			27	21	57	...	31	34		
D+201409+789+Like			26	41	29	...	7	35		
H+20140908+123+Like										
M+2014+123+Comment										

Store Graph in HBase

- Use property graph model to present a graph
- Use Hbase to store property graph model



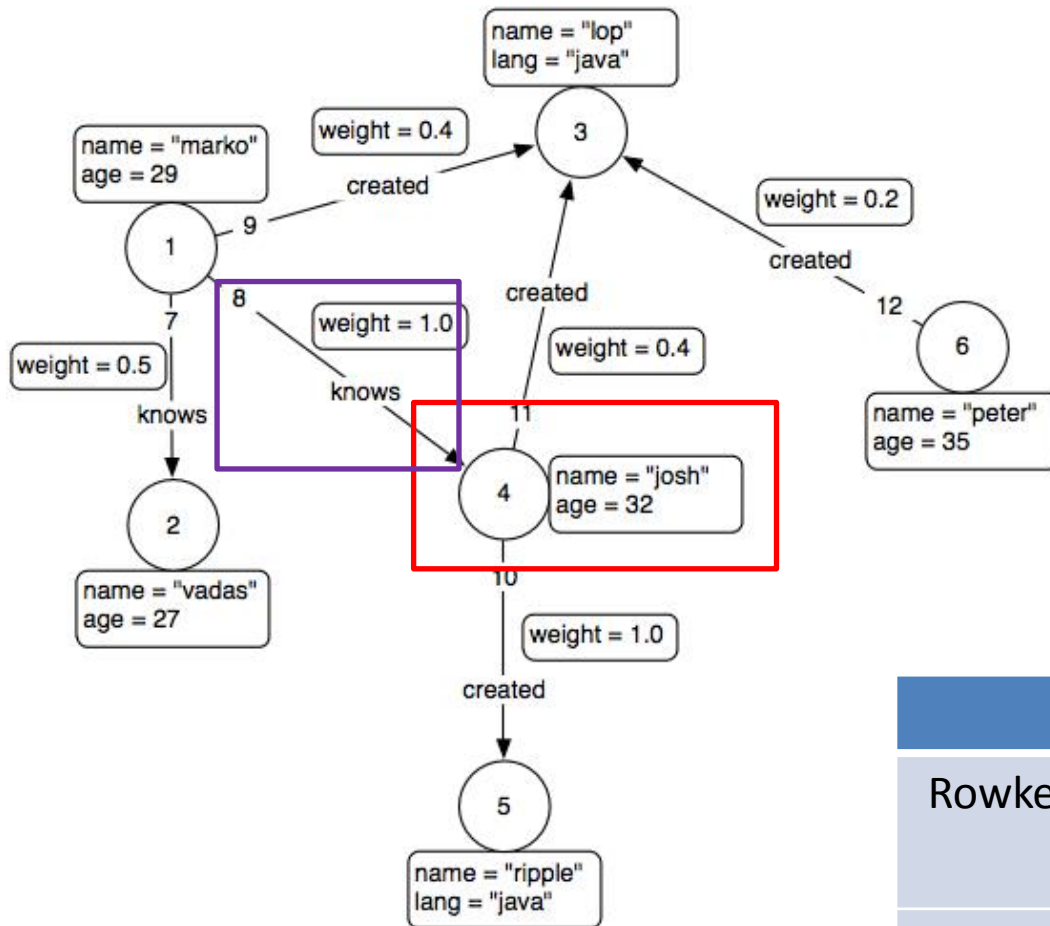


Vertex Table

rowkey vertex-id entity-type	Property (CF)	
	property-key1@property-value-type (qualifier)	...
	Property-value1	...

Edge Table

rowkey vertex1-row-key --> label --> vertex2-row-key	Property	
	property-key1@property-value-type	...
	Property-value1	...



Vertex Table

Rowkey	Property	
	name@String	age@int
4 MAN	josh	32

Edge Table

Rowkey	Property
	weight@float
1 WOMAN--> knows--> 4 MAN	1.0